DIFFERENTIABLE SATELLITE CONSTELLATION DESIGN VIA RELAXED COVERAGE AND REVISIT OBJECTIVES

Shreeyam Kacker
Department of Aeronautics and Astronautics
MIT
shreeyam@mit.edu

Kerri Cahoy
Department of Aeronautics and Astronautics
MIT
kcahoy@mit.edu

April 1, 2026

ABSTRACT

Keywords: satellite constellation design, differentiable programming, orbit propagation

1 Introduction

Designing satellite constellations is a fundamental problem in space mission design, across applications such as Earth Observation (EO), telecommunications, and navigation. The design space is high dimensional, with each spacecraft bringing up to six degrees of freedom corresponding to their orbital elements, and a wide range of often competing objectives to optimize over.

Traditional approaches address this complexity in one of two ways. Parametric families such as Walker Delta, Walker Star [21], and Flower constellations [15] reduce the design space to a fixed set of parameters but are inherently symmetric, making them non-optimal for non-uniform objectives. Metaheuristic methods such as genetic algorithms can handle arbitrary objectives but require thousands to millions of simulator rollouts [17].

For EO missions, the design problem is especially complex. Coverage and daytime revisit requirements may be non-uniform across ground targets, with additional constraints on data downlink. This non-uniformity makes parametric familie non-optimal, while the combinatorial nature of the problem makes metaheuristic methods computationally expensive.

Gradient-based optimization has also been previously attempted, but with limited success. Paek et al. [17] found the problem poorly conditioned for gradient-based methods, with difficulty even achieving convergence. Previous work has been limited the non-differentiability of the computational graph. Recent work by Acciarini et al. [3] has addressed part of this by developing ∂SGP4, an end-to-end differentiable SGP4 orbit propagator. However, the coverage and revisit metrics themselves remain non-differentiable, due to requiring binary visibility indicators, Boolean OR operations, and max functions.

To address this, we make the following contributions:

1. **Four continuous relaxations** that make coverage and revisit objectives differentiable: a soft sigmoid for visibility, a noisy-OR for multi-satellite aggregation, a leaky integrator for revisit gap tracking, and a LogSumExp soft-maximum for worst-case revisit (Figure 3).
2. **An end-to-end differentiable pipeline** from orbital elements to mission-level coverage metrics, composed with ∂SGP4 [3] and optimized with the AdamW optimizer (Figure 1).
3. **Demonstration on weighted target optimization**, where parametric families are inapplicable and prior gradient-based attempts failed.

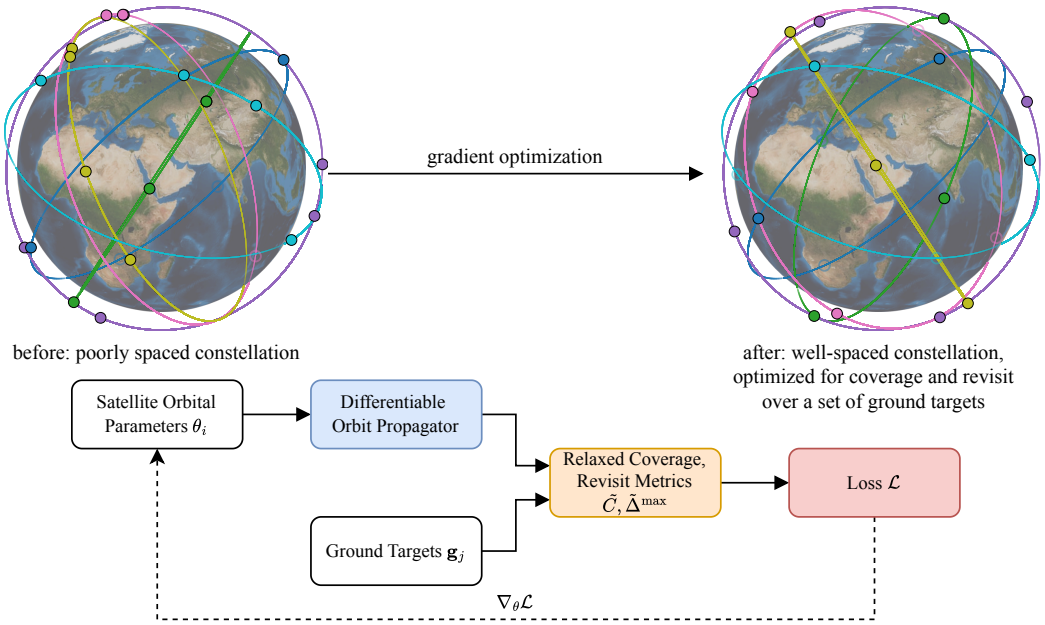

Figure 1: Overview of end-to-end differentiable constellation optimization. **Top:** Starting from a poorly-configured constellation with randomly spaced orbital planes and satellites, gradient optimization through the differentiable pipeline recovers a well-spaced constellation optimized for both coverage and revisit for a set of ground points. **Bottom:** Orbital parameters $\boldsymbol{\theta}_i$ are propagated through a differentiable orbit propagator and evaluated against a set of ground targets $\mathbf{g}_j$ to produce relaxed coverage and revisit metrics ($\tilde{C}$, $\tilde{\Delta}^{\max}$), which form a differentiable loss $\mathcal{L}$. Gradients $\nabla_{\boldsymbol{\theta}}\mathcal{L}$ flow back to the orbital parameters via reverse-mode automatic differentiation (dashed).

2 Related Work

2.1 Constellation Design and Optimization

Parametric constellations significantly reduce the number of design parameters, with the Walker Delta pattern [21] requiring only the number of planes, satellites per plane, and inclination. Flower constellations extend these ideas by using elliptical orbits to form specific 'harmonics' in revisit rates to produce repeating ground tracks [15], allowing for directed coverage over a specific area. These parametric families are efficient but inherently symmetric, and cannot be tailored to non-uniform or mission-specific coverage requirements without breaking the symmetry assumptions that make them tractable.

For non-parametric design, previous work has relied primarily on metaheuristic methods. Paek et al. [17] uses genetic algorithms and simulated annealing to reconfigure constellations, optimizing for global coverage and regional revisit over five design variables, requiring approximately 2500 total evaluations through expensive STK orbit simulations. They also attempted gradient-based optimization via finite differences and found it "extremely poorly conditioned," with convergence rarely achieved. In a survey of 144 papers on constellation design, Choo et al. [6] finds that gradient methods appear as a design approach in only two of the surveyed articles, making it among the least common methods—ranking far behind genetic algorithms, simulated annealing, and analytical approaches. The only gradient-based entry they identify, Fraire et al. [8], claims to be a gradient method, but in actuality uses a greedy neighbor search over discrete parameters such as number of satellites, planes, and inclination, with no analytic gradients computed. Abdelkhalik and Gad [2] uses a genetic algorithm to optimize single-orbit parameters for ground site coverage, explicitly noting that the fitness landscape contains numerous local minima that cause classical optimization methods to fail. Abdelkhalik and Gad [2] similarly claims to use a "second-order gradient" method, but no analytic gradients are used for optimization of design variables. Separately, Williams Rogers et al. [22] formulate constellation configuration as a collection of mixed-integer linear programs that select which satellites to activate from a pre-enumerated set of orbital slots, yielding provably optimal solutions for coverage and revisit objectives; however, the orbital slots and their underlying geometry (altitude, inclination, eccentricity) must be fixed a priori, so the method cannot discover improved orbit parameters.

Hwang et al. [11] demonstrates true gradient-based optimization of a satellite system through OpenMDAO, optimizing a single satellite's subsystem design across seven coupled disciplines with over 25,000 design variables, with gradient

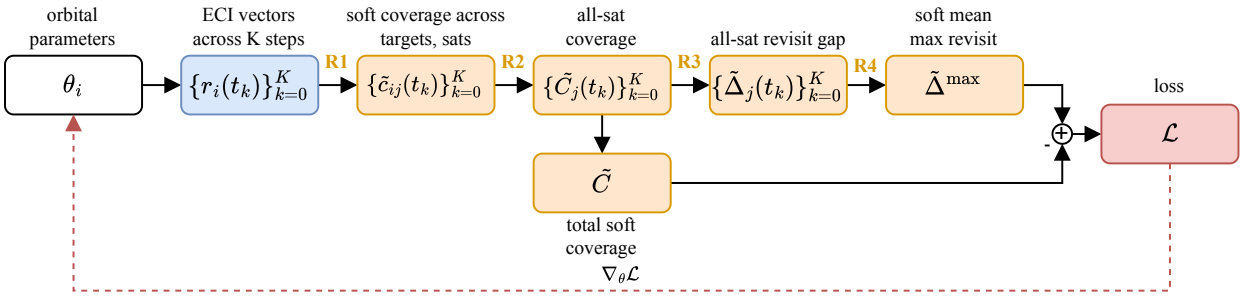

Figure 2: Computational graph for differentiable constellation optimization. Orbital parameters $\boldsymbol{\theta}_i$ are propagated to ECI position vectors $\{\mathbf{r}_i(t_k)\}$ across K timesteps. Four continuous relaxations (orange, **R1**–**R4**) transform the discrete coverage and revisit objectives into smooth surrogates: **R1** soft sigmoid visibility over targets and satellites, **R2** noisy-OR aggregation to all-satellite coverage, **R3** leaky integrator revisit gap accumulation, and **R4** LogSumExp soft-maximum over time. The coverage branch produces total soft coverage $\tilde{C}$; the revisit branch produces the soft mean worst-case revisit $\tilde{\Delta}^{\max}$. Both feed into the loss $\mathcal{L}$, and gradients $\nabla_{\boldsymbol{\theta}}\mathcal{L}$ flow back via reverse-mode automatic differentiation (dashed).

computation made tractable with adjoint state methods. Their work utilizes gradient-based objectives combined with orbit-dependent parameters, such as solar panel angles, battery sizing, communication scheduling, to optimize a single spacecraft's operation.

2.2 Differentiable Physics and Simulation

Differentiable programming [4] allows for gradients to be obtained across domains even where forward models are complex but smooth. The benefits of differentiable programming have been made significantly more accessible through frameworks that perform automatic differentiation, such as PyTorch and JAX. Hu et al. [10] demonstrated differentiable physical simulators targeting robotics and fluid dynamics, and Degrave et al. [7] develops differentiable physics engines to contact-rich robotics. For continuous relaxations applied to parameterized categorical distributions, Gumbel-Softmax reparameterization [12] is frequently used to compute analytic gradients. Kamra and Liu [13] proves a coverage gradient theorem for spatial multi-resource coverage objectives, providing an estimator for the gradient that can be used for optimization of sensor placement, via spatial discretization and implicit boundary differentiation applied to 2D planar sensor networks.

In astrodynamics, Acciarini et al. [3] introduces ∂SGP4, a PyTorch reimplementation of the SGP4/SDP4 orbit propagation algorithm, which is forward mode differentiable. Naylor et al. [16] extends this direction by combining a JAX-based differentiable SGP4 propagator with Mitsuba [1], a differentiable renderer, to optimize trajectories for spacecraft-to-spacecraft inspection operations, demonstrating end-to-end forward-mode differentiation through orbit propagation chained with a rendering pipeline.

This work builds on ∂SGP4 by introducing differentiable coverage and revisit objectives that enable direct optimization of mission performance metrics.

3 Method

3.1 Problem Formulation

We consider a constellation of N satellites, each parameterized by a two-line element (TLE) set. Each satellite brings up to a maximum of six degree of freedom corresponding to their orbital elements, $[a,\ e,\ i,\ \Omega,\ \omega,\ \mathcal{M}]^\top$, corresponding to semi-major axis, eccentricity, inclination, right ascension of the ascending node (RAAN), argument of periapsis, and mean anomaly, respectively. We universally fix the remaining TLE parameters corresponding to mean motion n, drag term B^*, and secular rates $\dot{n}$, $\ddot{n}$ to keep the constellation configuration static over time.

Given a grid of J ground target points $\mathcal{G} = \{\mathbf{g}_1, \ldots, \mathbf{g}_J\}$ on the Earth's surface and a time horizon $[0, T]$ discretized into K steps with spacing $\Delta t = T/(K-1)$, we seek to maximize coverage over the Earth and minimize the worst-case revisit gap.

Propagation. Each satellite's TLE elements are propagated through ∂SGP4 [3] to produce position vectors $\mathbf{r}_i(t_k)$ in the True Equator Mean Equinox (TEME) reference frame. These are rotated to the Earth-Centered Earth-Fixed (ECEF) frame via the Greenwich Mean Sidereal Time angle $\theta_{\text{GMST}}(t_k)$:

$$\mathbf{r}_i^{\text{ECEF}}(t_k) = \mathbf{R}_z\left(\theta_{\text{GMST}}(t_k)\right)\, \mathbf{r}_i^{\text{TEME}}(t_k), \tag{1}$$

where $\mathbf{R}_z(\theta)$ is the standard rotation matrix about the z-axis. This single-axis rotation is an approximation and neglects polar motion; the complete TEME-to-ECEF transformation requires an additional rotation between the Earth's instantaneous rotation axis and the ECEF pole [20]. For the $\sim 24\,\text{h}$ propagation horizons considered in this work, polar-motion offsets are on the order of tenths of arcseconds, corresponding to sub-kilometre position errors. These errors are small relative to the $\sim 2200\,\text{km}$ ground footprint of a satellite at $550\,\text{km}$ altitude with 10° minimum elevation.

Elevation angle. The elevation angle from ground point $\mathbf{g}_j$ to satellite i at time t_k is

$$\alpha_{ij}(t_k) = \arcsin\left(\frac{(\mathbf{r}_i^{\text{ECEF}}(t_k) - \mathbf{g}_j)\cdot \hat{\mathbf{g}}_j}{\|\mathbf{r}_i^{\text{ECEF}}(t_k) - \mathbf{g}_j\|}\right), \tag{2}$$

where $\hat{\mathbf{g}}_j = \mathbf{g}_j/\|\mathbf{g}_j\|$ is the local vertical unit vector at the ground point. A satellite is visible when $\alpha_{ij} \geq \alpha_{\min}$, where we consider the minimum elevation angle $\alpha_{\min} = 10^\circ$, which varies depending on mission requirements.

Coverage and revisit. Instantaneous coverage at ground point $\mathbf{g}_j$ at time t_k is the indicator function for when at least one satellite is visible:

$$C_j(t_k) = \mathbf{1}\left[\bigvee_{i=1}^{N} \alpha_{ij}(t_k) \geq \alpha_{\min}\right]. \tag{3}$$

The instantaneous coverage fraction over the full set of ground points $\mathcal{G}$ is then

$$C = \frac{1}{KJ}\sum_{k=1}^{K}\sum_{j=1}^{J} C_j(t_k). \tag{4}$$

The revisit gap at ground point $\mathbf{g}_j$ is the elapsed time since the most recent coverage event:

$$\Delta_j(t_k) = t_k - t_{k'}, \quad k' = \max\{k'' \leq k : C_j(t_{k''}) = 1\} \tag{5}$$

and the mean worst-case revisit across all ground points is then

$$\Delta^{\max} = \frac{1}{J}\sum_{j=1}^{J} \max_k \Delta_j(t_k). \tag{6}$$

We now introduce the relaxations required to make these metrics differentiable.

3.2 Relaxation 1: Soft Visibility

We replace the hard step function with a sigmoid (Figure 3a):

$$\tilde{c}_{ij}(t_k) = \sigma\left(\frac{\alpha_{ij}(t_k) - \alpha_{\min}}{\tau}\right), \quad \sigma(x) = \frac{1}{1+e^{-x}}, \tag{7}$$

where $\tau > 0$ controls the sharpness. As $\tau \to 0$ this recovers the hard indicator. This relaxation is analogous to the temperature-scaled softmax used in Gumbel-Softmax relaxations of discrete variables [12].

3.3 Relaxation 2: Noisy-OR Coverage Aggregation

Starting with the Boolean OR operator in Equation 3, treating each $\tilde{c}_{ij}(t_k)$ as an independent detection probability, the noisy-OR [19] relaxation gives (Figure 3b):

$$\tilde{C}_j(t_k) = 1 - \prod_{i=1}^{N}\left(1 - \tilde{c}_{ij}(t_k)\right). \tag{8}$$

This function is smooth, monotonically increasing in each $\tilde{c}_{ij}$, and recovers the hard OR when all inputs are in $\{0, 1\}$. This relaxation also inherently handles overlapping coverage: additional coverage of a point produces diminishing marginal gain, preventing the optimizer from clustering satellites.

The soft coverage fraction over the full set of ground points $\mathcal{G}$ with soft visibility and noisy-OR aggregation then becomes

$$\tilde{C} = \frac{1}{KJ} \sum_{k=1}^{K} \sum_{j=1}^{J} \tilde{C}_j(t_k). \tag{9}$$

3.4 Relaxation 3: Leaky Integrator for Revisit Gaps

The revisit gap at a ground point is the time elapsed since the most recent coverage event. Computing this gap requires identifying discrete transitions which we instead approximate with a leaky integrator recurrence [9] (Figure 3c):

$$\tilde{\Delta}_j(t_k) = \left(\tilde{\Delta}_j(t_{k-1}) + \delta t\right) \cdot \left(1 - \tilde{C}_j(t_k)\right), \quad \tilde{\Delta}_j(t_0) = 0. \tag{10}$$

When the ground point $\mathbf{g}_j$ is covered ($\tilde{C}_j \approx 1$), the accumulated gap is multiplied by approximately zero, resetting it. When the point is not covered ($\tilde{C}_j \approx 0$), the gap grows by Δt. This recurrence is fully differentiable and its maximum over time approximates the true worst-case revisit gap at each ground point.

3.5 Relaxation 4: LogSumExp Soft-Maximum

The soft worst-case revisit gap at ground point $\mathbf{g}_j$ is $\max_k \tilde{\Delta}_j(t_k)$, which has zero gradients almost everywhere. We replace it with the LogSumExp smooth maximum [5] (Figure 3d):

$$\widetilde{\Delta}_j^{\max} = \beta \cdot \log\left(\sum_{k=1}^{K} \exp\left(\frac{\tilde{\Delta}_j(t_k)}{\beta}\right)\right), \tag{11}$$

where β is a sharpness parameter, for which we assign $\beta = 30\,\text{min}$. This relaxation satisfies the inequality $\max_k \tilde{\Delta}_j(t_k) \leq \widetilde{\Delta}_j^{\max} \leq \max_k \tilde{\Delta}_j(t_k) + \beta \log K$, providing an upper bound. The mean worst-case revisit across all ground points is then

$$\tilde{\Delta}^{\max} = \frac{1}{J} \sum_{j=1}^{J} \widetilde{\Delta}_j^{\max}. \tag{12}$$

In practice, the coverage and revisit objectives have different sensitivity to the sigmoid temperature τ in Equation 7: phantom coverage is an additive bias for coverage but multiplicatively resets the revisit integrator. We therefore split τ into τ_{cov} for the coverage objective and τ_{rev} for the revisit computation, calibrated in §4.2.

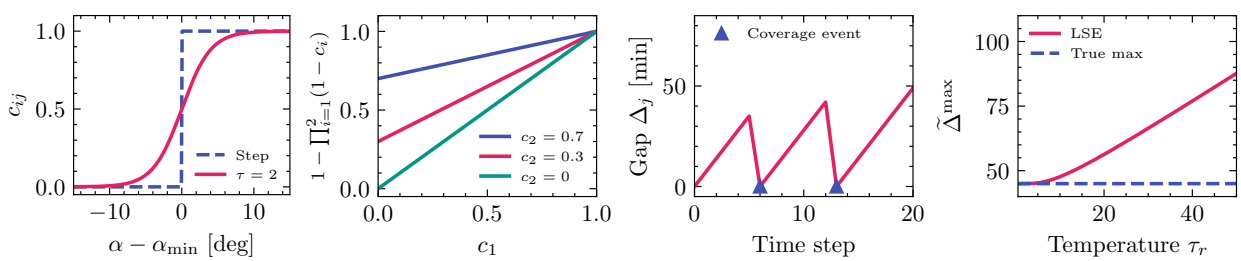

(a) Soft sigmoid (b) Noisy-OR (c) Leaky integrator (d) LogSumExp

Figure 3: The four continuous relaxations used in this work to produce a differentiable loss landscape.

3.6 Combined Objective

The total loss combines coverage maximization with revisit minimization:

$$\mathcal{L}(\mathbf{x}) = -\tilde{C} + \lambda\, \tilde{\Delta}^{\max}, \tag{13}$$

where $\mathbf{x}$ collects all free and reparameterized (§3.8) orbital parameters and λ is a weighting coefficient. The coverage term is instantaneous: it penalizes overlapping fields of view and saturates once satellites are well-separated. The revisit term, aggregated over the full simulation, continues providing gradients after coverage plateaus. Although both terms

push toward spacing satellites apart, they operate at different scales—coverage quickly separates clusters, while revisit works globally.

For objectives that are not geodesically uniform, the sums over ground points can be replaced with weighted sums over a discrete target set (§4.6).

3.7 Tightness of the Composed Relaxation

Each of the four relaxations is individually tight: the soft sigmoid recovers the hard step as $\tau \to 0$; the noisy-OR is exact when its inputs are binary; the leaky integrator is exact when coverage is in $\{0, 1\}$; and the LogSumExp satisfies $\max_k x_k \leq \text{LSE} \leq \max_k x_k + \beta \log K$. However, the *composition* of these relaxations does not inherit these guarantees, because each relaxation receives the soft output of the previous one rather than the binary values for which it was designed.

The interaction introduces two competing biases:

Phantom coverage (gap shortening). The soft sigmoid coverage indicator assigns nonzero coverage $\tilde{c}_{ij} > 0$ to satellites slightly *below* the visibility threshold ($\alpha_{ij} < \alpha_{\min}$). This phantom coverage propagates through the noisy-OR and partially resets the leaky integrator, producing gaps that are *shorter* than the true discrete gaps.

Incomplete reset (gap lengthening). Conversely, a satellite well above the threshold produces $\tilde{c}_{ij} \approx 0.99$, not exactly 1. The leaky integrator retains a fraction $(1 - \tilde{C}_j) \approx 0.01$ of the accumulated gap at each coverage event. Over time, this residual prevents the gap from fully resetting, producing relaxed gaps that are *longer* than the true gaps during sustained coverage. The LogSumExp adds a further upward bias of at most $\beta \log K$.

These effects compete: phantom coverage shortens gaps near the visibility boundary, while incomplete resets and the LogSumExp overestimate lengthen them. The net bias depends on the specific geometry and is neither a guaranteed upper nor lower bound on the true objective. This asymmetric sensitivity—additive for coverage, multiplicative for revisit—motivates the separated sigmoid temperatures discussed in §3.5.

The relaxation temperatures control a tradeoff: smaller values tighten the approximation closer to the true hard metrics, but sharpen the gradients, potentially destabilizing optimization. We investigate this tradeoff empirically in §4.2 by computing the hard coverage and revisit metrics at the optimizer's solution and comparing them to the relaxed values, as well as through PCA over the optimizer landscape (Figure 6).

3.8 Optimization

Computational graph. The full forward pass composes the ∂SGP4 propagator with the four relaxations into a single differentiable graph (Figure 2). Every operation is implemented in PyTorch [18], so $\nabla_{\theta_i} \mathcal{L}$ is computed via reverse-mode automatic differentiation through the entire graph. The Jacobian $\partial \mathbf{r}_i^{\text{TEME}} / \partial \mathbf{x}_i$ is handled inside ∂SGP4: it is dense and time-varying—a small change in RAAN rotates the entire orbital plane, affecting position at all future time steps.

Interval constraints and reparameterization. Without constraints, the optimizer trivially raises altitude until all ground points are permanently visible. We enforce interval constraints (e.g. LEO altitude bounds) via sigmoid reparameterization, mapping unconstrained variables θ_i to valid elements $\mathbf{x}_i$:

$$x_d = \ell_d + (u_d - \ell_d)\, \sigma(\theta_d) \quad \Rightarrow \quad \frac{\partial x_d}{\partial \theta_d} = (u_d - \ell_d)\, \sigma(\theta_d)(1 - \sigma(\theta_d)), \tag{14}$$

for bounded parameters (e.g. inclination) with bounds $[\ell_d, u_d]$, and $x_d = \theta_d$ (identity) for periodic parameters (RAAN, mean anomaly) where wrapping is natural. This reparameterization eliminates the need for projected gradient or constrained optimization methods. For coupled constraints—such as requiring a minimum perigee altitude regardless of eccentricity—the per-element sigmoid is insufficient since perigee $r_p = a(1 - e)$ couples semi-major axis and eccentricity. We reparameterize in (perigee altitude, excess altitude) space:

$$r_p = r_p^{\min} + (r_p^{\max} - r_p^{\min})\, \sigma(\theta_{r_p}), \tag{15}$$

$$\delta r = \delta r^{\max}\, \sigma(\theta_{\delta r}), \tag{16}$$

$$a = R_\oplus + r_p + \tfrac{1}{2}\delta r, \quad e = \frac{\delta r}{2a}, \tag{17}$$

where r_p is the perigee altitude, $\delta r = r_a - r_p \geq 0$ is the excess altitude (apogee minus perigee), and both are independently sigmoid-bounded. This construction guarantees $r_p \geq r_p^{\min}$ and $e \geq 0$ by construction, while allowing the optimizer to freely explore eccentric orbits.

Plane constraints via gradient accumulation. Conventionally, satellite constellations are designed with multiple satellites sharing a single plane, to allow for launching multiple satellites per launch vehicle. Rather than enforcing these per-plane constraints as hard constraints, we average the gradients per-plane: if satellites in a plane share a single RAAN parameter, the gradient is the average of per-satellite contributions. The averaged gradient is applied to the shared parameter, and the updated value is synchronized across all satellites in the plane after each optimizer step. This reduces the effective dimensionality without modifying the forward pass, and generalizes to any shared parameters: per-plane inclination sharing, altitude coupling across the constellation, or any grouping governed by mission constraints.

4 Experiments

We first calibrate the relaxation hyperparameters to ensure the relaxed loss is representative of true solutions. We then evaluate the framework on three experiments: a toy problem validating gradient correctness, a uniform-coverage baseline that recovers known-optimal geometries, and a weighted target optimization on a problem class where parametric families fail.

4.1 Common Setup

All experiments use ∂SGP4 [3] for orbit propagation with PyTorch autograd for gradient computation. Circular orbits are assumed and held stationary by zeroing the mean-motion gradient. The minimum elevation angle for all ground points is $\alpha_{\min} = 10^\circ$. As described in §4.2, we use separate sigmoid temperatures for coverage and revisit: the coverage temperature τ_{cov} controls gradient smoothness, while a tighter revisit temperature τ_{rev} prevents the leaky integrator from being gamed by phantom coverage. Coverage and revisit metrics are evaluated over a $24\,\text{h}$ propagation horizon discretized into $K = 240$ time steps on a ground grid of $36 \times 72 = 2592$ points spanning $\pm 70^\circ$ latitude. Specific relaxation parameters, optimizer settings, and revisit weight λ vary per experiment and are stated in each subsection.

4.2 Hyperparameter Tuning

The composed relaxation has four hyperparameters that define the loss surface geometry: the coverage sigmoid temperature τ_{cov}, the revisit sigmoid temperature τ_{rev} (introduced in §3.2), the LogSumExp temperature β, and the revisit weight λ. We exclude optimizer dynamics (learning rate, momentum) from this calibration, as they affect convergence speed but not the shape of the loss landscape. We calibrate by running four optimizations of a 24-satellite constellation (6×4, inclination fixed at 60°, RAAN and mean anomaly free, per-plane RAAN sharing) from different initial RAAN configurations:

1. **Near-uniform**: $\Omega = [0, 60, 120, 180, 240, 300]^\circ$ (close to even spacing),
2. **Moderate**: $\Omega = [0, 30, 120, 200, 210, 300]^\circ$ (irregular but spread),
3. **Clustered**: $\Omega = [0, 10, 20, 30, 40, 50]^\circ$ (all planes within 50°),
4. **Two-cluster**: $\Omega = [0, 5, 180, 185, 270, 275]^\circ$ (three tight clusters).

All tuning runs use the same random mean anomaly initialization and default relaxation parameters ($\tau = 2^\circ$, $\beta = 10\,\text{min}$, $\lambda = 0.1$) for 800 AdamW iterations.

Ranking failure with shared softness. The near-uniform and moderate initializations converge to solutions with hard mean revisit gaps of $\sim 48\,\text{min}$, while the clustered and two-cluster initializations find solutions with gaps of $\sim 68\,\text{min}$. However, under the default relaxation parameters, the relaxed loss of the $68\,\text{min}$ solutions is *lower* than that of the $48\,\text{min}$ solutions—the relaxation ranks the worse solutions as better.

This ranking failure is caused by the shared sigmoid temperature: the same τ controls both the coverage objective and the revisit leaky integrator (Equation 10). A soft sigmoid with $\tau = 2^\circ$ assigns coverage $c_{ij} \approx 0.27$ to a satellite at $\alpha = \alpha_{\min} - 2^\circ$. For the coverage fraction (Equation 9), this phantom coverage has minimal effect because it is a small additive bias. But in the leaky integrator, it *multiplicatively* resets the gap: $\Delta_j \leftarrow \Delta_j \cdot (1 - C_j)$. When multiple clustered satellites each contribute partial coverage at low elevation, the noisy-OR combines them into substantial phantom coverage ($C_j \gg 0$) that partially resets the gap timer at every time step. The clustered solutions benefit disproportionately from this effect.

Split-softness resolution. We decouple the sigmoid temperature into τ_{cov} for the coverage objective and τ_{rev} for the revisit integrator, then grid-search over $\tau_{\text{cov}} \in \{1.0, 1.5, 2.0, 2.5, 3.0, 4.0, 5.0\}^\circ$, $\tau_{\text{rev}} \in$

$\{0.1, 0.25, 0.5, 0.75, 1.0, 1.5, 2.0\}^\circ$, $\beta \in \{3, 5, 7.5, 10, 15, 20\}\mathrm{min}$, and $\lambda \in \{0.05, 0.1, 0.2, 0.5, 1.0, 2.0\}$ (1764 combinations, each requiring only four gradient-free forward passes). A combination is *valid* if it correctly ranks the good solutions (near-uniform and moderate) below the bad solutions (clustered and two-cluster).

Table 1 summarizes the findings. No valid combinations exist when $\tau_{\text{cov}} = \tau_{\text{rev}} \geq 2^\circ$ (the original shared-softness regime). Decoupling the temperatures resolves the ranking failure: combinations with $\tau_{\text{cov}} \geq 3^\circ$ and $\tau_{\text{rev}} \leq 0.5^\circ$ correctly rank solutions while maintaining smooth coverage gradients. See Appendix A for the full grid search results.

Smoothness–correctness tradeoff. The valid hyperparameter combinations tend toward tight revisit sigmoids ($\tau_{\text{rev}} \leq 1^\circ$), which produce correct global rankings but sharp, poorly conditioned gradients. In practice, optimization converges more reliably with softer values (e.g., $\tau_{\text{cov}} = 5^\circ$, $\tau_{\text{rev}} = 2^\circ$) that lie outside the valid regime. This is not contradictory: the calibration tests *global* ranking correctness across qualitatively different solutions, but gradient-based optimization only requires *local* gradient quality within a basin of attraction. Softer relaxations smooth the loss surface, improving convergence, while the underlying basin structure still guides the optimizer toward good solutions. The hyperparameters used in subsequent experiments reflect this tradeoff, favoring reliable convergence over exact global ranking.

Table 1: Split-softness calibration results. Valid = relaxed loss correctly ranks solutions with better hard metrics below those with worse hard metrics. Decoupling the coverage and revisit sigmoid temperatures allows smooth coverage gradients with a tight revisit sigmoid.

Regime	**Softness range**	**Valid combos**
Shared $\tau_{\text{cov}} = \tau_{\text{rev}} \geq 2^\circ$	Original setting	0
Shared $\tau_{\text{cov}} = \tau_{\text{rev}} = 1^\circ$	Tighter, sharp gradients	Some (with $\beta \leq 5$, $\lambda \geq 0.2$)
Split: $\tau_{\text{cov}} \geq 2^\circ$, $\tau_{\text{rev}} \leq 0.5^\circ$	Smooth + tight	Many
Split: $\tau_{\text{cov}} \geq 3^\circ$, $\tau_{\text{rev}} \leq 1^\circ$	Recommended	Many

4.3 Experiment 1: Toy Problem

To validate that the relaxed objectives produce correct gradients and results, we consider the simplest non-trivial constellation design problem: two satellites in a single orbital plane at $550\,\mathrm{km}$ altitude with inclination fixed at 60° and RAAN fixed at 0°, optimizing only their mean anomalies $\mathcal{M}_1$ and $\mathcal{M}_2$. We use relaxation parameters $\tau_{\text{cov}} = 5^\circ$, $\tau_{\text{rev}} = 2^\circ$, $\beta = 10\,\mathrm{min}$, $\lambda = 2.0$, and AdamW with learning rate $\eta = 10^{-2}$.

The optimal solution is known analytically: the two satellites should be separated by 180° in mean anomaly to maximize coverage and revisit rate within the shared orbital plane. This reduces the problem to two free parameters, allowing the full loss landscape to be visualized as a 2D contour plot.

Loss landscape. We evaluate both the relaxed loss and hard (discrete) metrics on a dense grid over $(\mathcal{M}_1, \mathcal{M}_2) \in [0, 360)^{2\circ}$. Figure 4 shows the resulting landscapes: the top row displays the relaxed loss, coverage, and revisit, while the bottom row shows the corresponding hard metrics. Both landscapes exhibit a clear global minimum along the anti-diagonal $|\mathcal{M}_1 - \mathcal{M}_2| = 180^\circ$, with no local minima. The relaxed landscape is visibly smoother than the hard landscape while preserving the same basin structure, confirming that the relaxations provide a smooth surface amenable to gradient descent without distorting the underlying objective.

Gradient descent trajectory. Starting from $\mathcal{M}_1 = 120^\circ$, $\mathcal{M}_2 = 130^\circ$ (a 10° separation), Adam converges to $|\mathcal{M}_1 - \mathcal{M}_2| \approx 180^\circ$ within 400 iterations. The trajectory (overlaid in white on Figure 4) follows the gradient field smoothly toward the minimum without oscillation. This validates that the composed relaxations produce well-behaved gradients—a prerequisite for the higher-dimensional experiments that follow.

4.4 Experiment 2: Uniform Coverage Baseline

Before applying the method to non-uniform objectives, we establish a baseline on the well-understood uniform coverage problem. This validates that the optimizer can recover known-optimal geometries and provides a point of comparison for the weighted-target experiment.

We optimize $N = 24$ satellites in 6 orbital planes (4 per plane) with inclination fixed at 60°. RAAN and mean anomaly are free, with RAAN shared within each plane, yielding $6 + 24 = 30$ effective degrees of freedom. We use AdamW with learning rate $\eta = 10^{-2}$ and revisit weight $\lambda = 0.1$.

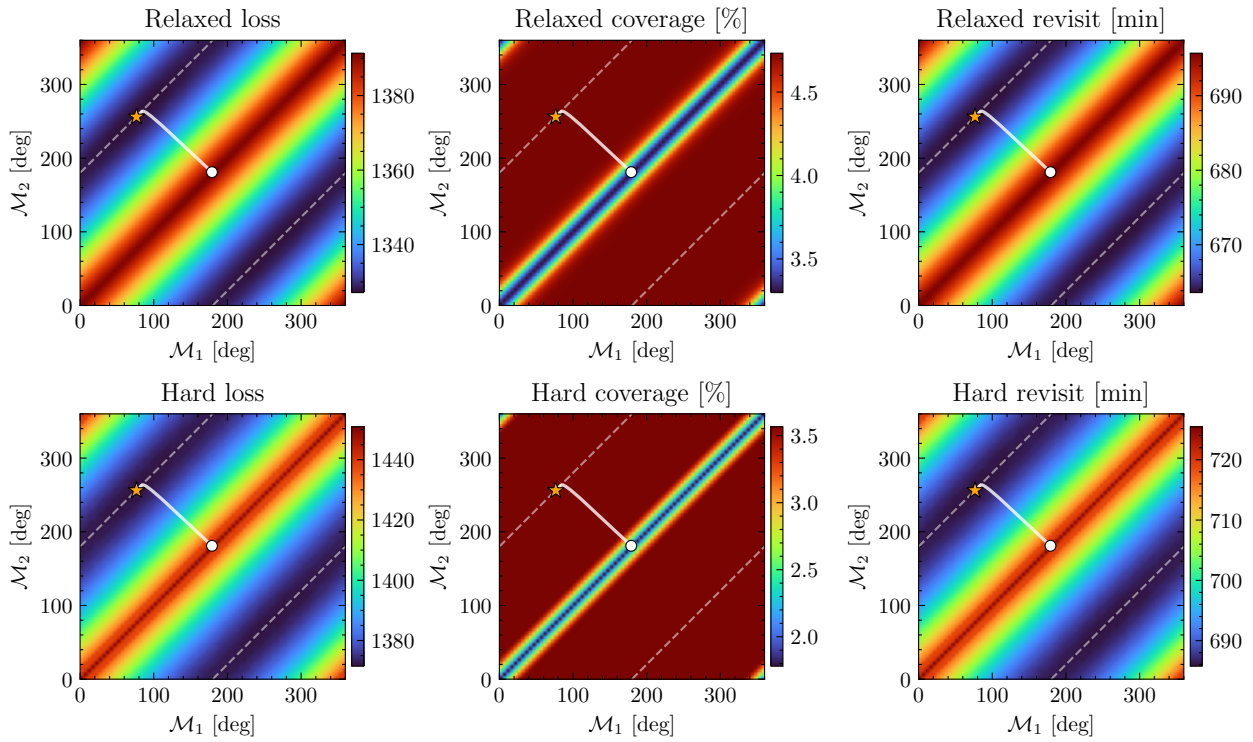

Figure 4: Loss landscape for the two-satellite mean anomaly optimization. **Top row:** relaxed loss, coverage, and revisit. **Bottom row:** corresponding hard (discrete) metrics. The global minimum lies along $|\mathcal{M}_1 - \mathcal{M}_2| = 180°$ (dashed line). The gradient descent trajectory (white line, circle = start, star = end) converges from a $10°$ separation to the optimum. The relaxed landscape is smoother than the hard landscape while preserving the same basin structure.

Initial configuration. We deliberately initialize the constellation with irregular RAAN offsets $[0, 30, 120, 200, 210, 300]°$ (far from uniform $60°$ spacing) and random mean anomalies (seed 42). The classical Walker-delta $24/6/1$ pattern uses $60°$ RAAN spacing.

Results. Over 1000 iterations, the optimizer recovers near-uniform RAAN spacing from the irregular initialization, approaching the Walker-delta geometry without domain-specific heuristics. Table 2 summarizes the metrics; Figure 5 shows convergence, spatial coverage redistribution, and ground tracks.

Table 2: Walker recovery: 24-satellite constellation (6×4), 1000 iterations, $550\,\mathrm{km}$ fixed altitude, inclination fixed at $60°$. Soft metrics are the relaxed surrogate values; hard metrics are the discrete ground truth.

Metric	**Initial**	**Optimized**	**Walker 24/6/1**
Hard coverage [%]	32.41	40.34	40.33
Hard mean max revisit [min]	136.2	48.0	48.0
Soft coverage [%]	41.86	41.96	50.32
Soft mean max revisit [min]	127.5	73.5	73.5

4.5 Loss Landscape Analysis

Following Li et al. [14], we visualize the loss landscape by computing per-trajectory principal components and evaluating the landscape on each trajectory's own 2D slice through the 30-dimensional parameter space. Unlike a single global PCA projection—which discards the dimensions that matter for individual trajectories—each slice faithfully represents the landscape that its optimizer actually navigated. Figure 6 shows smooth, well-conditioned basins across all four initializations from §4.2, in contrast to the "extremely poorly conditioned" landscape reported by Paek et al. [17]

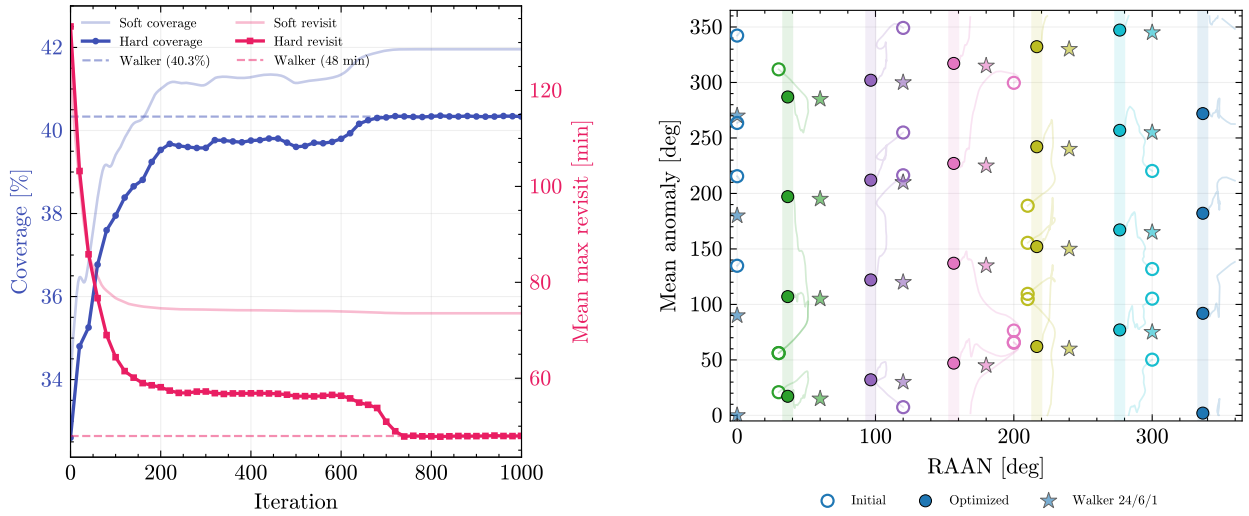

(a) Convergence

(b) RAAN vs mean anomaly

Figure 5: Walker recovery experiment. **(a)** Convergence of soft and hard coverage and revisit metrics over 1000 iterations, with Walker 24/6/1 reference (dashed). The optimizer matches Walker performance by iteration ∼400. **(b)** RAAN vs mean anomaly trajectories. Hollow circles show the initial configuration, solid circles the optimized result, and stars the Walker reference.

when using finite-difference gradients through discrete objectives. This confirms that the poor conditioning observed in prior work was a consequence of the non-differentiable objective formulation, not the underlying physics.

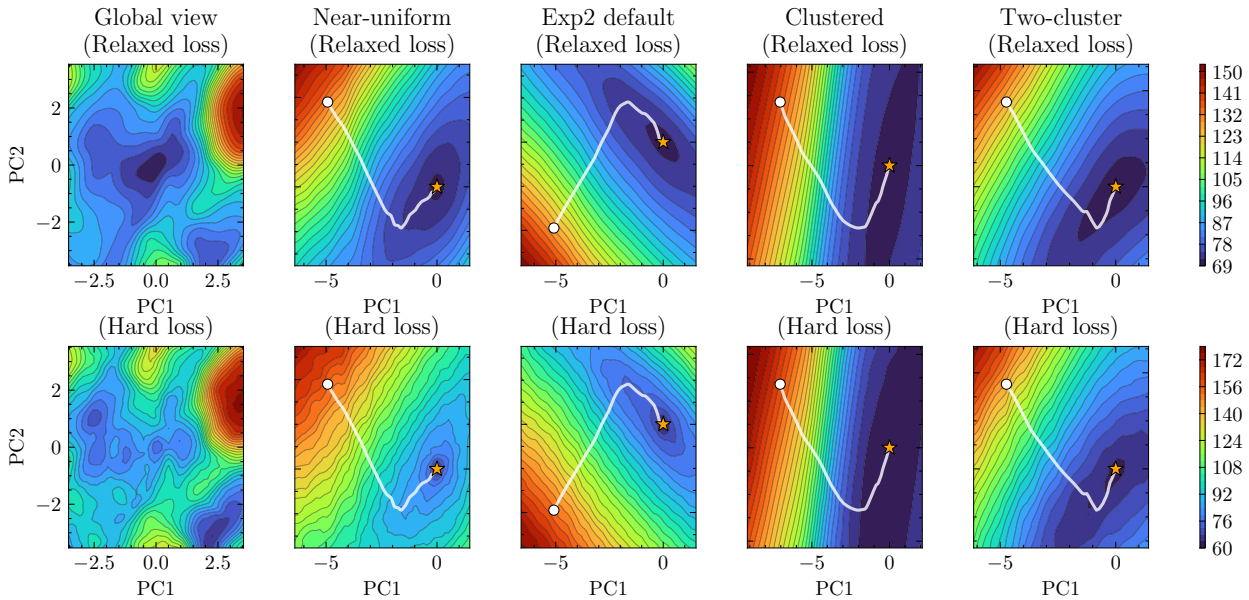

Figure 6: Loss landscape visualization [14]. **Left column:** global view along random directions (zoomed out) revealing the broader landscape structure. **Columns 2–5:** per-trajectory PCA slices, each showing the landscape along that trajectory's own principal components (white line, circle = start, star = end). **Top row:** relaxed total loss. **Bottom row:** hard (discrete) total loss. The per-trajectory views show smooth basins, while the global view reveals the multi-modal structure of the full parameter space.

4.6 Experiment 3: Regional Target Optimization

The primary advantage of gradient-based constellation optimization over parametric families is the ability to handle non-uniform, mission-specific objectives. We demonstrate this by optimizing a small constellation for coverage and revisit over Europe—a problem that Walker constellations cannot address by construction.

Setup. We sample 500 ground targets from European country boundaries (GeoJSON), weighted by $\cos(\text{lat})$ to correct for area distortion. The constellation consists of 4 satellites in 2 orbital planes (2 per plane). Unlike Experiment 2, we free all orbital geometry parameters: inclination ($\iota \in [30, 90]^\circ$), RAAN, mean anomaly, argument of perigee, and orbital shape. To allow the optimizer to explore eccentric orbits without producing unphysical perigee altitudes, we reparameterize eccentricity and altitude jointly as (perigee altitude, excess altitude), where perigee $\in [400, 600]\,\text{km}$ and excess $\in [0, 1500]\,\text{km}$, guaranteeing perigee $\geq 400\,\text{km}$ by construction. Inclination, RAAN, eccentricity, argument of perigee, and altitude are shared within each plane; mean anomaly is free per satellite, giving 12 effective degrees of freedom. We use AdamW with $\eta = 10^{-2}$, $\lambda = 1.0$, and run for 2000 iterations.

Results. Starting from a Walker-like circular initialization at $550\,\text{km}$, the optimizer discovers a qualitatively different geometry (Table 3). Both planes converge to highly elliptical orbits ($e \approx 0.40$) with argument of perigee $\omega \approx 270^\circ$, placing apogee over the northern hemisphere where the satellite dwells longest. The resulting orbits have perigee $\approx 685\,\text{km}$ and apogee $\approx 10\,000\,\text{km}$—a Molniya-class highly elliptical orbit discovered entirely from gradients, with no domain-specific heuristics. The inclination converges to $\sim$63°–65°, near the critical inclination (63.4°) at which J_2 perturbations do not precess the argument of perigee.

The optimized constellation achieves 99.3% coverage and 3.7 min mean revisit over Europe, compared to 10.0% coverage and 247.5 min for a Walker 4/2/1 baseline (Table 3).

Figure 8 shows the visibility density (number of timesteps each ground bin is within 10° elevation of any satellite). The optimized constellation concentrates nearly all coverage in the northern hemisphere over Europe, while the Walker baseline distributes coverage uniformly but achieves far less density. The orbital period of $\sim 210\,\text{min}$ produces approximately 7 revolutions per sidereal day, creating a near-resonant repeating ground track.

Table 3: Regional target optimization: 4 satellites (2 planes $\times$ 2), 3000 iterations, all orbital geometry parameters free with coupled perigee/apoapsis constraints.

Metric	**Walker 4/2/1**	**Optimized**
Hard coverage [%]	9.97	99.25
Hard mean max revisit [min]	247.5	3.7
Soft coverage [%]	12.37	97.99
Soft mean max revisit [min]	208.9	54.9

Table 4: Per-plane orbital geometry discovered by the optimizer for the Europe target set. Both planes converge to eccentric, Molniya-like orbits with apogee over the northern hemisphere ($\omega \approx 270^\circ$). The Walker baseline uses circular orbits at $550\,\text{km}$.

	ι [deg]	Ω [deg]	e	ω [deg]	Perigee / apogee [km]
Plane 0	62.6	336.6	0.398	276.2	685 / 10011
Plane 1	65.3	155.6	0.402	266.5	687 / 10193
Walker (all)	60.0	uniform	0.001	—	550 / 550

5 Conclusions and Future Work

In conclusion, we show that gradient-based methods when composed with the appropriate relaxations and propagation techniques, are viable for design of satellite constellations.

The relaxed metrics introduced in this work track hard metrics (§4.2) and produce well-conditioned gradients, in contrast to previous attempts using finite-difference approaches [17].

On uniform coverage objectives with plane-accumulation constraints, we can derive Walker-Delta constellations from purely our differentiable objectives.

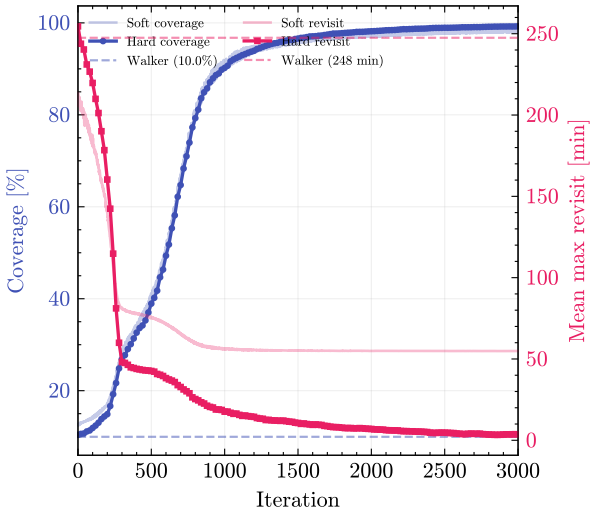

Figure 7: Convergence of the Europe target optimization over 3000 iterations. The optimizer rapidly improves both coverage and revisit over the Walker baseline (dashed).

On weighted regional targets, the optimizer can discover asymmetric geometries, such as eccentric orbits that dwell over target latitudes, without requiring any hand-crafted inductive biases.

Limitations. The current approach optimizes a single-epoch snapshot and does not account for long-term perturbations (J2 RAAN drift, atmospheric drag) that alter constellation geometry over longer time horizons. While excluded in this work by zeroing out relevant gradients, these factors can be included with appropriate metrics to evaluate correctness, for example, total spacecraft lifetime above a certain altitude in the case of atmospheric drag. In this work, we discretize time steps, which introduces approximation error in the revisit metrics, requiring significant propagation over a fixed time horizon to acquire gradients. This can introduce noise into the gradients and require more compute, although still likely much less than for a different optimization method, like an evolutionary algorithm. While the contributions in this work allow for differentiable objectives, design variables for the constellation design problem (number of planes, satellites per plane) remain non-differentiable and must be set by the user or searched over separately. In practice,

Future work. Incorporating J2 secular perturbations would enable optimization of long-term constellation maintenance, potentially discovering geometries that exploit RAAN drift for improved coverage evolution. While this work shows one application of differentiable objectives for space mission design, the applications even in constellation design are numerous. For example, investigating optimal placement of satellites given hetereogenous instruments with different field of view, optimal satellite constellation reconfiguration for new objectives, satellite constellation design with differently weighted ground points with different minimum elevations, and more.

References

[1] Mitsuba 3 - A Retargetable Forward and Inverse Renderer. URL `https://www.mitsuba-renderer.org/`.

[2] Ossama Abdelkhalik and Ahmed Gad. Optimization of space orbits design for Earth orbiting missions. *Acta Astronautica*, 68(7-8):1307–1317, April 2011. ISSN 00945765. doi:10.1016/j.actaastro.2010.09.029. URL `https://linkinghub.elsevier.com/retrieve/pii/S0094576510003796`.

[3] Giacomo Acciarini, Atılım Güneş Baydin, and Dario Izzo. Closing the gap between SGP4 and high-precision propagation via differentiable programming. *Acta Astronautica*, 226:694–701, January 2025. ISSN 00945765. doi:10.1016/j.actaastro.2024.10.063. URL `https://linkinghub.elsevier.com/retrieve/pii/S0094576524006374`.

[4] Atılım Gunes Baydin, Barak A Pearlmutter, Alexey Andreyevich Radul, and Jeffrey Mark Siskind. Automatic Differentiation in Machine Learning: a Survey. 2018.

[5] Stephen P. Boyd and Lieven Vandenberghe. *Convex optimization*. Cambridge University Press, Cambridge New York Melbourne New Delhi Singapore, version 29 edition, 2023. ISBN 978-0-521-83378-3.

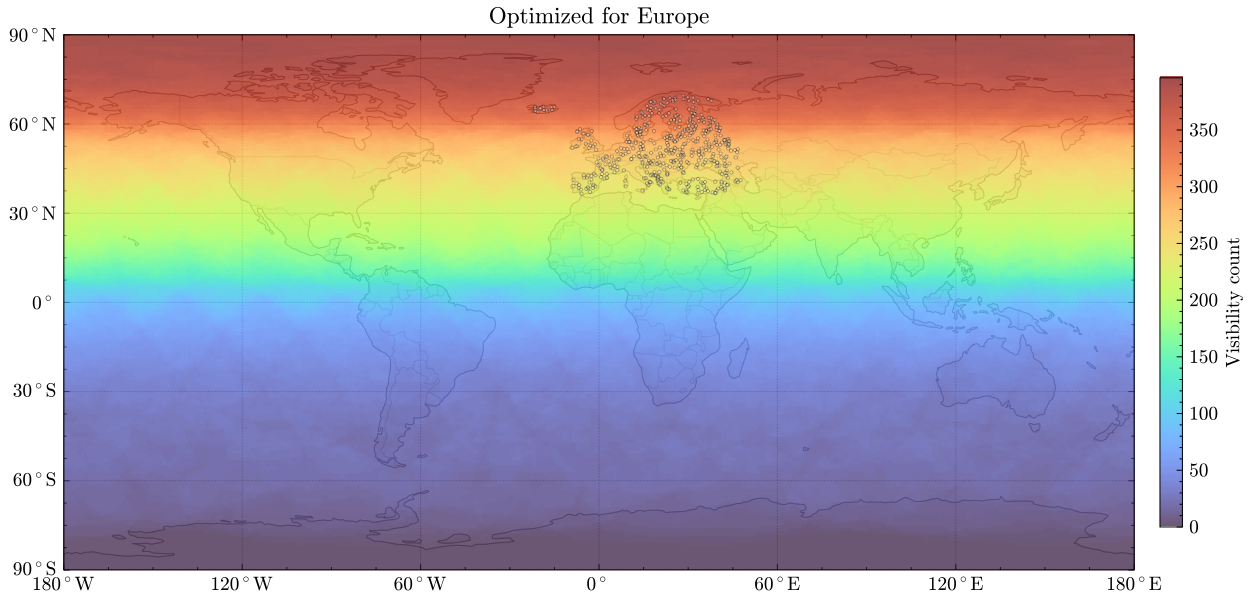

(a) Optimized for Europe

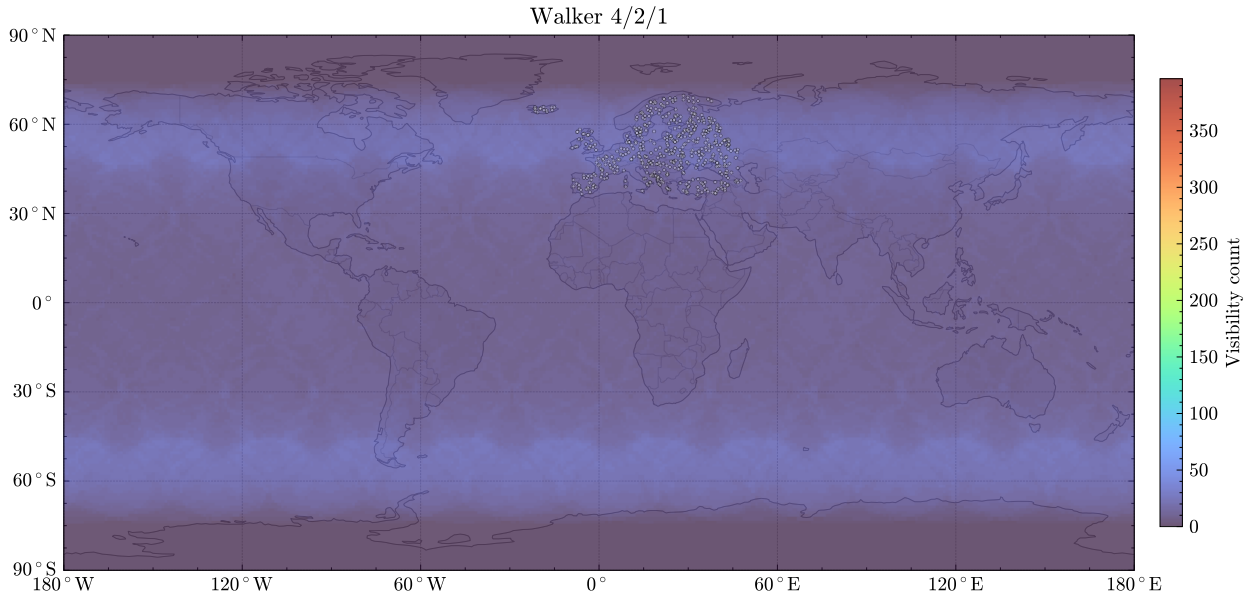

(b) Walker 4/2/1 baseline

Figure 8: Visibility density over a $24\,\mathrm{h}$ window: number of timesteps each ground bin is within 10° elevation of any satellite. **(a)** The optimized constellation concentrates nearly all coverage over European latitudes via highly elliptical orbits ($e \approx 0.4$, $\omega \approx 270^\circ$), with apogee dwell over the northern hemisphere. **(b)** The Walker baseline distributes coverage symmetrically but achieves far less density over Europe.

[6] Nathaniel Choo, Darryl Ahner, and Bryan Little. A Survey of Orbit Design and Selection Methodologies. *The Journal of the Astronautical Sciences*, 71(1):4, January 2024. ISSN 2195-0571. doi:10.1007/s40295-023-00420-9. URL `https://link.springer.com/10.1007/s40295-023-00420-9`.

[7] Jonas Degrave, Michiel Hermans, Joni Dambre, and Francis Wyffels. A Differentiable Physics Engine for Deep Learning in Robotics. *Frontiers in Neurorobotics*, 13:6, March 2019. ISSN 1662-5218. doi:10.3389/fnbot.2019.00006. URL `https://www.frontiersin.org/article/10.3389/fnbot.2019.00006/full`.

[8] Juan Fraire, Santiago Henn, Fabio Dovis, Roberto Garello, and Giorgio Taricco. *Sparse Satellite Constellation Design for LoRa-based Direct-to-Satellite Internet of Things*. December 2020. doi:10.1109/GLOBECOM42002.2020.9348042. Pages: 6.

[9] Wulfram Gerstner and Werner M. Kistler. *Spiking Neuron Models: Single Neurons, Populations, Plasticity*. Cambridge University Press, August 2002. ISBN 978-0-521-89079-3. Google-Books-ID: Rs4oc7HfxIUC.

[10] Yuanming Hu, Luke Anderson, Tzu-Mao Li, Qi Sun, Nathan Carr, Jonathan Ragan-Kelley, and Frédo Durand. DiffTaichi: Differentiable Programming for Physical Simulation, February 2020. URL `http://arxiv.org/abs/1910.00935`. arXiv:1910.00935 [cs].

[11] John T. Hwang, Dae Young Lee, James W. Cutler, and Joaquim R. R. A. Martins. Large-Scale Multidisciplinary Optimization of a Small Satellite's Design and Operation. *Journal of Spacecraft and Rockets*, 51(5):1648–1663, September 2014. ISSN 0022-4650, 1533-6794. doi:10.2514/1.A32751. URL `https://arc.aiaa.org/doi/10.2514/1.A32751`.

[12] Eric Jang, Shixiang Gu, and Ben Poole. Categorical Reparameterization with Gumbel-Softmax, August 2017. URL `http://arxiv.org/abs/1611.01144`. arXiv:1611.01144 [cs, stat].

[13] Nitin Kamra and Yan Liu. Gradient-based optimization for multi-resource spatial coverage problems. In *Proceedings of the Thirty-Seventh Conference on Uncertainty in Artificial Intelligence*, pages 1885–1894. PMLR, December 2021. URL `https://proceedings.mlr.press/v161/kamra21a.html`.

[14] Hao Li, Zheng Xu, Gavin Taylor, Christoph Studer, and Tom Goldstein. Visualizing the Loss Landscape of Neural Nets. In *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc., 2018. URL `https://proceedings.neurips.cc/paper_files/paper/2018/hash/a41b3bb3e6b050b6c9067c67f663b915-Abstract.html`.

[15] Daniele Mortari, Matthew P. Wilkins, and Christian Bruccoleri. The Flower Constellations. *The Journal of the Astronautical Sciences*, 52(1-2):107–127, March 2004. ISSN 0021-9142, 2195-0571. doi:10.1007/BF03546424. URL `https://link.springer.com/10.1007/BF03546424`.

[16] Jack Naylor, Raghav Mishra, Nicholas H. Barbara, and Donald G. Dansereau. dLITE: Differentiable Lighting-Informed Trajectory Evaluation for On-Orbit Inspection, 2025. URL `https://arxiv.org/abs/2512.16011`. Version Number: 1.

[17] Sung Wook Paek, Sangtae Kim, and Olivier de Weck. Optimization of Reconfigurable Satellite Constellations Using Simulated Annealing and Genetic Algorithm. *Sensors*, 19(4):765, January 2019. ISSN 1424-8220. doi:10.3390/s19040765. URL `https://www.mdpi.com/1424-8220/19/4/765`. Number: 4.

[18] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Kopf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. PyTorch: An Imperative Style, High-Performance Deep Learning Library. In *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019. URL `https://proceedings.neurips.cc/paper_files/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html`.

[19] Judea Pearl. *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. Elsevier, June 2014. ISBN 978-0-08-051489-5. Google-Books-ID: mn2jBQAAQBAJ.

[20] David Vallado, Paul Crawford, Ricahrd Hujsak, and T. S. Kelso. Revisiting Spacetrack Report #3. In *AIAA/AAS Astrodynamics Specialist Conference and Exhibit*. American Institute of Aeronautics and Astronautics. doi:10.2514/6.2006-6753. URL `https://arc.aiaa.org/doi/abs/10.2514/6.2006-6753`. _eprint: https://arc.aiaa.org/doi/pdf/10.2514/6.2006-6753.

[21] J. G. Walker. Satellite Constellations. *Journal of the British Interplanetary Society*, 37:559, December 1984. ISSN 0007-084X. URL `https://ui.adsabs.harvard.edu/abs/1984JBIS...37..559W`. ADS Bibcode: 1984JBIS...37..559W.

[22] David O. Williams Rogers, Dongshik Won, Dongwook Koh, Kyungwoo Hong, and Hang Woon Lee. Optimal Satellite Constellation Configuration Design: A Collection of Mixed Integer Linear Programs. *Journal of Spacecraft and Rockets*, pages 1–18, March 2026. ISSN 0022-4650. doi:10.2514/1.A36518. URL `https://arc.aiaa.org/doi/10.2514/1.A36518`.

A Hyperparameter Tuning

This appendix provides the complete experimental setup and raw results for hyperparameter tuning described in §4.2.

Constellation configuration. The tuning uses a 24-satellite constellation (6×4) at $550\,\mathrm{km}$ altitude with eccentricity ≈ 0, inclination fixed at $60°$, and RAAN and mean anomaly as free parameters. RAAN is shared within each plane, giving 30 effective degrees of freedom (6 unique RAANs + 24 mean anomalies).

Optimization settings. Each of the four initial configurations is optimized for 800 iterations with AdamW ($\eta = 10^{-2}$), propagation over $24\,\mathrm{h}$ discretized into $K = 240$ time steps, a ground grid of $36 \times 72 = 2592$ points spanning $\pm 70°$ latitude with $\cos(\mathrm{lat})$ area weights, minimum elevation $\alpha_{\min} = 10°$, sigmoid temperature $\tau = 2°$, LogSumExp temperature $\beta = 10\,\mathrm{min}$, revisit weight $\lambda = 0.1$, mean revisit reduction, and GMST randomization enabled. The optimization uses the *default* relaxation parameters, not the calibrated ones—the goal is to obtain four diverse final solutions whose relative quality is known from the hard metrics.

Initial RAAN configurations. The four configurations are: near-uniform $[0, 60, 120, 180, 240, 300]°$, moderate $[0, 30, 120, 200, 210, 300]°$, clustered $[0, 10, 20, 30, 40, 50]°$, and two-cluster $[0, 5, 180, 185, 270, 275]°$. All use the same random mean anomaly initialization (NumPy `RandomState(42)`, uniform on $[0, 360)°$).

Grid search parameters. The grid search evaluates the relaxed loss of all four final solutions under each combination of: coverage softness $\tau_{\text{cov}} \in \{1.0, 1.5, 2.0, 2.5, 3.0, 4.0, 5.0\}°$, revisit softness $\tau_{\text{rev}} \in \{0.1, 0.25, 0.5, 0.75, 1.0, 1.5, 2.0\}°$, LogSumExp temperature $\beta \in \{3.0, 5.0, 7.5, 10.0, 15.0, 20.0\}\mathrm{min}$, and revisit weight $\lambda \in \{0.05, 0.1, 0.2, 0.5, 1.0, 2.0\}$. This gives $7 \times 7 \times 6 \times 6 = 1764$ combinations, each requiring 4 gradient-free forward passes (one per solution), for 7056 evaluations total.

Validity criterion. A combination is *valid* if the relaxed loss correctly ranks the good solutions below the bad:

$$\max\big(\mathcal{L}_{\text{near-uniform}},\, \mathcal{L}_{\text{moderate}}\big) < \min\big(\mathcal{L}_{\text{clustered}},\, \mathcal{L}_{\text{two-cluster}}\big). \tag{18}$$

The *margin* is the difference between the right- and left-hand sides. Among valid combinations, we recommend the one that maximizes τ_{cov} (smoothest coverage gradients), then τ_{rev}, with margin as tiebreaker.

Reproducibility. The experiment is implemented in `experiments/exp0_hyperparam_tuning.py`. The full list of valid combinations is written to `experiments/exp0_valid_combos.txt`, sorted by descending coverage softness.